

A Summary of the Current State of the Art Algorithms for the Asymmetric Traveling Salesman Problem

Alex Wong (asw2181) - asw2181@columbia.edu

May 1, 2017

Abstract

We look at the state of the art algorithm for the general case of the Asymmetric Traveling Salesman problem and its current known inapproximability bounds. The currently best algorithm, by Asadpour et al.[2], details an $O(\log n / \log \log n)$ -approximation algorithm for the general case of ATSP for costs satisfying the triangle inequality. An algorithm by Karpinski et al.[10] gives us a hardness of approximation bounds of $75 / 74$ for ATSP. We also briefly touch on some recent work on approximation factors of special cases of ATSP, including an approach by Svensson et al.[14] that details a constant factor approximation algorithm for a special case of ATSP that contains only two different edge weights of arbitrary weight where the costs satisfy the triangle inequality.

1 Introduction

The traveling salesman problem is one of the most well-known and studied problems in theoretical computer science and operations research. Many real world problems, from delivery routing to latency networks to DNA sequencing, can be framed as some variation of the traveling salesman problem. In this paper, we look at the current state of the art algorithms for the asymmetric traveling salesman problem (ATSP) where the costs satisfy the triangle inequality (costs from some point u to point w is at most the cost from point u to point v + cost from point v to point w) but the cost from some point u to point v does not necessarily equal the cost from point v to point u . From this description, we can frame the symmetric traveling salesman problem (the cost from some point u to point v does equal the cost from point v to point u) as a special case of ATSP.

In Section 2, we provide an overview of an $O(\log n / \log \log n)$ -approximation algorithm for the general case of ATSP for costs satisfying the triangle inequality which comes from the work of Asadpour et al. [2]. The $O(\log n / \log \log n)$ -approximation algorithm introduces the concept of the "thinness" of a spanning tree where the thinness correlates with the approximation bound of ATSP. This is an improvement over the decades-long standing $\Theta(\log n)$ -approximation bound given by the work of Frieze et al. [5].

In section 4, we detail some of the recent work for special cases of ATSP. Gharan and Saberi showed that there is a constant factor approximation on ATSP of planar graphs and graphs with bounded genus [6]. Anari and Gharan showed that the integrality gap of ATSP is $O(\text{poly log log } n)$ [1], but there is no current method known to construct a solution of $O(\text{poly log log } n)$ -approximation. The recent work of Svensson et al. showcased a new heuristic for tackling approximating ATSP [13], with a constant factor approximation on a special case of ATSP that contains only two different edge weights [14].

Lastly, in section 5, we conclude with opinions on the hypothesized next steps towards a constant factor approximation algorithm for ATSP.

2 An $O(\log n / \log \log n)$ -approximation algorithm for ATSP

Algorithm 1 details the $O(\log n / \log \log n)$ -approximation algorithm for ATSP (from [2] section 1).

Algorithm 1 An $O(\log n / \log \log n)$ -approximation algorithm for ATSP [2]

Input: A set V consisting of n points and a cost function $c : V \times V \rightarrow \mathbb{R}^+$ where the costs satisfy the triangle inequality.

Output: An $O(\frac{\log n}{\log \log n})$ -approximation to the ATSP instance described by the inputs V and c .

1. Solve the Held-Karp LP relaxation of the ATSP instance to get an optimum extreme point solution \mathbf{x}^* . Create a symmetrized and scaled down version of \mathbf{x}^* from (5) and define it as \mathbf{z}^* . \mathbf{z}^* is a vector which can be interpreted as a point in the relative interior (not touching any boundaries) of the spanning tree polytope P created from an undirected graph supported by \mathbf{x}^* where we disregard the directions of arcs. Another way to think of \mathbf{z}^* is the marginal probabilities on the edges z_e^* of an exponential distribution on spanning trees $\tilde{p}(\cdot)$.

2. Let E be the support graph of \mathbf{z}^* where we disregard the directions of arcs. We then find weights $\{\tilde{\gamma}\}_{e \in E}$ such that $\tilde{p}(T)$ is approximately proportional to $\exp(\sum_{e \in T} \tilde{\gamma}_e)$ where, for any edge $e \in E$,

$$\sum_{T \in \mathcal{T} : T \ni e} \tilde{p}(T) \leq (1 + \epsilon) z_e^*, \text{ for a small value } \epsilon.$$

3. Sample $\Theta(\log n)$ spanning trees T_i from $\tilde{p}(\cdot)$. Orient all the edges of the sampled spanning trees to minimize its cost since the spanning tree only has to be weakly connected. Let \vec{T} represent an oriented tree. Let T^* be the spanning tree that has minimum total cost amongst all the sample spanning trees.

4. Find a minimum cost integral circulation that contains the oriented tree \vec{T}^* . This will give us a multigraph. Shortcut the multigraph to get a Eulerian tour and output the result.

Throughout the rest of this section, we will divide each step of the algorithm into subsection(s) to describe it in more detail and/or prove its correctness.

2.1 Preliminaries and Notation

In this subsection, we describe much of the notation that will be used throughout the rest of this section. We will also define the notion of "thinness" of a spanning tree.

Let $a = (u, v)$ be the arc (directed edge) from u to v and let $e = (u, v)$ be an undirected edge. Then, A (resp. E) is the set of arcs (resp. edges) in a directed (resp. undirected) graph G .

For a given function $f : A \rightarrow \mathbb{R}$, the cost of f is defined as $c(f) := \sum_{a \in A} c(a)f(a)$. For some set of arcs $S \subseteq A$, we define $f(S) := \sum_{a \in S} f(a)$. The same notation is used for the edge set E of an undirected graph.

For a subset of vertices $U \subseteq V$, we define

$$\delta^+(U) := \{a = (u, v) \in A : u \in U, v \notin U\},$$

$$\delta^-(U) := \{a = (u, v) \in A : u \notin U, v \in U\},$$

$$A(U) := \{a = (u, v) \in A : u \in U, v \in U\}$$

$$\delta(U) := \delta^+(U) \cup \delta^-(U),$$

to be the set of arcs that are leaving, entering, contained in U , and total of incoming and outgoing arcs, respectively. We also define $\delta^+(v) := \delta^+(\{v\})$ and $\delta^-(v) := \delta^-(\{v\})$ for each single vertex v . For

an undirected graph, $\delta(U)$ represents the set of edges with just one endpoint in U and $E(U)$ represents the edges that are contained within U . Lastly, all \log in equations represents the natural logarithm.

We say that a spanning tree T is α -thin with respect to \mathbf{z} if and only if:

$$|T \cap \delta(U)| \leq \alpha \cdot z(\delta(U)) \quad \forall U \subset V$$

which essentially states that for all the subsets of vertices U of the spanning tree T , the number of edges in the cut between U and $V \setminus U$ must be \leq the thinness multiplied by the cost of the edges in the cut.

We also say that T is (α, s) -thin with respect to \mathbf{z} if and only if it is α -thin and $c(T) \leq s \cdot OPT_{HK}$, which essentially states that the cost of T is at most s times the cost of the optimal Held-Karp solution.

By showing that a spanning tree is "thin", we would be able to get a Eulerian augmentation of the spanning tree where the total cost is within a factor of α of the cost OPT_{HK} , which implies within the same factor of the optimum solution.

2.2 The Held-Karp Relaxation

We detail and prove step 1 of the algorithm. Given an ATSP instance with a cost function $c : V \times V \rightarrow \mathbb{R}^+$, we can get a lower bound on the optimal value of ATSP by the following linear programming (LP) relaxation defined on the complete bidirected graph over the vertex set V :

$$\min \sum_a c(a) x_a \tag{1}$$

$$s.t. \quad \mathbf{x}(\delta^+(U)) \geq 1 \quad \forall U \subset V, \tag{2}$$

$$\mathbf{x}(\delta^+(v)) = \mathbf{x}(\delta^-(v)) = 1 \quad \forall v \in V, \tag{3}$$

$$x_a \geq 0 \quad \forall a.$$

This relaxation is accurate since (2) makes sure that all the vertices are strongly connected and (3) makes sure that the indegree = outdegree for all vertices of V which makes this a Eulerian graph. This relaxation is known as the Held-Karp relaxation [9] and it's well-known that an optimal solution \mathbf{x}^* to the relaxation can be computed in polynomial time (using the ellipsoid method). Thus, we can say that $c(\mathbf{x}^*) = OPT_{HK}$. Also, notice that (3) implies that any feasible solution \mathbf{x} satisfies

$$\mathbf{x}(\delta^+(U)) = \mathbf{x}(\delta^-(U)) \quad \forall U \subset V. \tag{4}$$

2.3 Proving $\mathbf{z}^* \in \text{Relative Interior of } P$

As mentioned in the algorithm, \mathbf{z}^* is the symmetrized and scaled down version of \mathbf{x}^* . So, we formally define

$$\mathbf{z}_{\{u,v\}}^* := \frac{n-1}{n} (\mathbf{x}_{uv}^* + \mathbf{x}_{vu}^*). \tag{5}$$

As mentioned in step 2 of the algorithm, we let E be the support of \mathbf{z}^* . We also let A be the support of \mathbf{x}^* where $A = \{(u,v) : x_{uv}^* > 0\}$. We also define the costs of all edges $e \in E$ where $c(e) = \min\{c(a) : a \in \{(u,v), (v,u)\} \cap A\}$, $\forall e \in E$. This shows that $c(\mathbf{z}^*) < c(\mathbf{x}^*)$. Now we can prove the rest of step 1 of the algorithm and define a lemma:

Lemma 2.1 ([2], section 3): The vector \mathbf{z}^* belongs to the relative interior of the spanning tree polytope P .

Proof: From the characterization of the base polytope of a matroid given by Edmonds[4], it follows that the spanning tree polytope P is defined by the following inequalities (see Corollary 50.7c of [12]):

$$P = \{z \in \mathbb{R}^E : z(E) = |V| - 1, \tag{6}$$

$$z(E(U)) \leq |U| - 1 \quad \forall U \subset V \text{ where } U \neq \emptyset, \quad (7)$$

$$z_e \geq 0 \quad \forall e \in E. \quad (8)$$

Thus, for some vector z to belong in the relative interior of P , it must satisfy both inequalities (7) and (8) strictly ($z(E(U)) < |U| - 1$ and $z_e > 0$).

From the above equations, we see that \mathbf{z}^* satisfies constraint (6) because:

$$\forall v \in V, \mathbf{x}^*(\delta^+(v)) = 1 \Rightarrow \mathbf{x}^*(A) = n = |V| \Rightarrow \mathbf{z}^*(E) = n - 1 = |V| - 1$$

Now, consider any set $U \subset V$ where $U \neq \emptyset$. We get that

$$\begin{aligned} \sum_{v \in U} \mathbf{x}^*(\delta^+(v)) &= |U| = \mathbf{x}^*(A(U)) + \mathbf{x}^*(\delta^+(U)) \geq \mathbf{x}^*(A(U)) + 1 \\ &\Rightarrow \mathbf{x}^*(A(U)) \leq |U| - 1 \end{aligned}$$

We see that \mathbf{x}^* satisfies constraints (2) and (3), so we can use equation (5) to show

$$\begin{aligned} \mathbf{z}^*(E(U)) &= \frac{n-1}{n} \mathbf{x}^*(A(U)) < \mathbf{x}^*(A(U)) \leq |U| - 1 \\ &\Rightarrow \mathbf{z}^*(E(U)) < |U| - 1 \end{aligned}$$

where we see that \mathbf{z}^* satisfies inequality (7) strictly. Also, since E is the support of \mathbf{z}^* , that means \mathbf{z}^* also satisfies inequality (8) strictly. Thus, we can conclude that \mathbf{z}^* is in the relative interior of P . \square

We also make a note that since \mathbf{x}^* is an optimum extreme point solution, it's known that its support A has at most $3n - 4$ arcs (see [8], Theorem 15). Also, \mathbf{x}^* can be expressed as the unique solution (due to global minima of a convex program) of an invertible system with only 0 - 1 coefficients, then every entry \mathbf{x}_a^* is rational with an integral numerator and denominator bounded by $2^{O(n \log n)}$, thus we can say that $\mathbf{z}_{min}^* > 2^{-O(n \log n)}$ (see [2] section 3).

2.4 Thin Trees and Their Relation to ATSP

A crucial element to the algorithm of approximating ATSP is the notion of thin trees. In section 2.1, we briefly described what "thinness" means and what properties a thin tree has. We will now detail it further to showcase why thin trees are crucial to this approximation algorithm and prove its correctness.

We jump to step 4 of the algorithm to prove it. First, though, a theorem that is required for the proof of correctness of thin trees is Hoffman's circulation theorem. A circulation is essentially any function $f : A \rightarrow \mathbb{R}$ such that $f(\delta^+(v)) = f(\delta^-(v))$ for each vertex $v \in V$. Hoffman's circulation theorem gives a necessary and sufficient condition for the existence of circulation on the arcs subject to the arcs' lower and upper capacities. The following is Hoffman's circulation theorem (we won't prove Hoffman's circulation theorem; see [12], Theorem 11.2 for the proof):

Theorem 2.2 (Hoffman's Circulation Theorem[12]): Given lower and upper capacities $l, u : A \rightarrow \mathbb{R}$, there exists a circulation f satisfying $l(a) \leq f(a) \leq u(a)$ for all $a \in A$ if and only if:

1. $l(a) \leq u(a)$ for all $a \in A$ and
2. for all subsets $U \subseteq V$, we have $l(\delta^-(U)) \leq u(\delta^+(U))$.

Furthermore, if l and u are integer-valued, f can be chosen to be integer-valued too.

Given this theorem on circulation, we now provide a theorem that proves that the ability to find an (α, s) -thin spanning tree with respect to \mathbf{z}^* , for some α and s , translates directly into an ability to obtain an $(2\alpha + s)$ -approximation to ATSP:

Theorem 2.3 ([2], section 4): Let \mathbf{x}^* be an optimal solution to the Held-Karp relaxation and \mathbf{z}^* to be the symmetrized and scaled down version of \mathbf{x}^* defined in (5). if T^* is an (α, s) -thin spanning tree with respect to \mathbf{z}^* for some α and s , then we can find, in polynomial-time, a Hamiltonian cycle whose cost is at most $(2\alpha + s)c(\mathbf{x}^*) = (2\alpha + s)OPT_{HK} \leq (2\alpha + s)OPT$.

Proof: We first orient the edges (u, v) of T^* to minimize the cost of T^* where each arc $a \in T^* = \min\{c(a) : a \in \{(u, v), (v, u)\} \cap A\}$, and we denote the oriented tree as \vec{T}^* . Notice that by the definition of the undirected cost function that we defined back in section 2.3, we have $c(\vec{T}^*) = c(T^*)$.

Lets now consider an augmentation of minimum cost to make \vec{T}^* an Eulerian directed graph. To get the augmentation, we can translate it as a minimum cost circulation problem with integral lower capacities and a bound on the upper capacities bsd on what we want to prove here in Theorem 2.3. So, we first set the lower capacities of the arcs by

$$l(a) = \begin{cases} 1 & a \in \vec{T}^* \\ 0 & a \notin \vec{T}^* \end{cases}$$

and the minimum circulation problem is $\min\{c(f) : f \text{ is a circulation and } f(a) \geq l(a), \forall a \in A\}$. This minimum circulation problem is well known so we do not prove it here (see [12], Corollary 12.2a for the proof). The minimum cost circulation problem's proof will give us an optimal integral circulation solution f^* in polynomial time. When we augment this solution with \vec{T}^* , we end up with a directed multigraph that is Eulerian (indegree = outdegree $\forall v \in V$), which we will define as H . Thus, H is strongly connected so we can take a Eulerian tour and shortcut it to get a Hamiltonian cycle of cost $\leq c(f^*)$.

The last part of this proof is showing an upper bound on $c(f^*)$. As we stated in our theorem, we are looking for cost $\leq (2\alpha + s)c(\mathbf{x}^*)$, meaning that $c(f^*) \leq (2\alpha + s)c(\mathbf{x}^*)$. So, we can define the upper capacities of the arc by

$$u(a) = \begin{cases} 1 + 2\alpha x_a^* & a \in \vec{T}^* \\ 2\alpha x_a^* & a \notin \vec{T}^*. \end{cases}$$

Now, we claim that there \exists a circulation g where $l(a) \leq g(a) \leq u(a) \forall a \in A$. This implies that

$$c(f^*) \leq c(g) \leq c(u) = c(\vec{T}^*) + 2\alpha c(\mathbf{x}^*) \leq (2\alpha + s)c(\mathbf{x}^*),$$

which gives us the bound we want for $c(f^*)$. To prove this claim, we see that the α -thinness of T^* implies that for any $U \subseteq V$, the number of arcs of $\vec{T}^* \in \delta^-(U) \leq \alpha \mathbf{z}^*(\delta(U))$, regardless of the orientation where T^* oriented to \vec{T}^* . Thus, we can say that

$$l(\delta^-(U)) \leq \alpha \mathbf{z}^*(\delta(U))$$

and by using a combination of (4) and (5), we can say that

$$l(\delta^-(U)) \leq \alpha \mathbf{z}^*(\delta(U)) < 2\alpha \mathbf{x}^*(\delta^-(U)).$$

We also know that

$$2\alpha \mathbf{x}^*(\delta^+(U)) \leq u(\delta^+(U)).$$

By equation (4), we can say that

$$2\alpha \mathbf{x}^*(\delta^+(U)) = 2\alpha \mathbf{x}^*(\delta^-(U)).$$

So, we can finally conclude

$$\begin{aligned} l(\delta^-(U)) &\leq 2\alpha \mathbf{x}^*(\delta^-(U)) = 2\alpha \mathbf{x}^*(\delta^+(U)) \leq u(\delta^+(U)) \\ &\Rightarrow l(\delta^-(U)) \leq u(\delta^+(U)) \end{aligned}$$

for any $U \subseteq V$, thus showing that the circulation g does exist, which concludes the proof of this theorem. \square

2.5 Maximum Entropy Distribution

In this subsection, we detail the approach on how to find an entropy distribution $p(\cdot)$. We let \mathcal{T} represent the collection of all spanning trees of $G = (V, E)$ and let \mathbf{z} represent some arbitrary point in the relative interior of the spanning tree polytope P of G . We now let $p^*(\cdot)$ be the maximum entropy distribution with respect to the marginal probabilities imposed by \mathbf{z} . This is considered the optimum solution of the following convex program:

$$\begin{aligned} \inf \quad & \sum_{T \in \mathcal{T}} p(T) \log p(T) \\ \text{s.t.} \quad & \sum_{T \ni e} p(T) = z_e \quad \forall e \in E, \\ & p(T) \geq 0 \quad \forall T \in \mathcal{T}. \end{aligned} \tag{9}$$

This convex program is feasible since \mathbf{z} is in the relative interior of the spanning tree polytope P and as the objective function is bounded and the feasible region is compact, the infimum (greatest lower bound) is found, which conveys the existence of an optimal solution $p^*(\cdot)$. Also, since the objective function is strictly convex (any local minima is essentially a global minima), this implies that the maximum entropy distribution $p^*(\cdot)$ is unique. OPT_{CP} will represent the optimum value of the convex program (9).

We define $p(T)$ as the probability of sampling a spanning tree T in $p(\cdot)$. Also, for any feasible solution $p(\cdot)$, we have $\sum_T p(T) = 1$ because

$$n - 1 = \sum_{e \in E} z_e = \sum_{e \in E} \sum_{T \ni e} p(T) = (n - 1) \sum_T p(T)$$

based on the constraints given by convex program (9).

Now, assuming that \mathbf{z} is in the relative interior of the spanning tree polytope P , we want to now show that $p^*(T) > 0 \quad \forall T \in \mathcal{T}$ and that $p^*(T)$ gives a simple exponential formula. To show this, we find the Lagrange to the convex program (9). Thus, $\forall e \in E$, we have a Lagrange multiplier δ_e (not to be confused with δ^- and δ^+ which represents incoming and outgoing edges, respectively) to the constraint corresponding to the marginal probability z_e , and we define the Lagrange function as

$$L(p, \delta) = \sum_{T \in \mathcal{T}} p(T) \log p(T) - \sum_{e \in E} \delta_e \left(\sum_{T \ni e} p(T) - z_e \right)$$

which can be rewritten as

$$L(p, \delta) = \sum_{e \in E} \delta_e z_e + \sum_{T \in \mathcal{T}} \left(p(T) \log p(T) - p(T) \sum_{e \in T} \delta_e \right)$$

thus making the Lagrange dual to the convex program (9)

$$\sup_{\delta} \inf_{p \geq 0} L(p, \delta) \tag{10}$$

We solve the inner infimum. Since each of the contributions of $p(T)$ is separable, then $\forall T \in \mathcal{T}$, $p(T)$ must minimize the convex function $p(T) \log p(T) - p(T) \delta(T)$ where $\delta(T) = \sum_{e \in T} \delta_e$. Now, taking partial derivatives of the convex function with respect to $p(T)$, we get

$$1 + \log p(T) - \delta(T) = 0$$

$$\log p(T) = \delta(T) - 1$$

$$p(T) = e^{\delta(T)-1} \tag{11}$$

So, the inner infimum

$$\inf_{p \geq 0} L(p, \delta) = \sum_{e \in E} \delta_e z_e - \sum_{T \in \mathcal{T}} e^{\delta(T)-1}$$

In the latter part of step 2 of the algorithm, to find the weights $\tilde{\gamma}_e, \forall e \in E$, requires lengthy non-trivial calculations which will not be detailed within this paper. There are two methods to find $\tilde{\gamma}_e$'s: either the combinatorial approach (see [2] section 7) where the basic idea is to iteratively tweak the weights of $\tilde{\gamma}_e$ until their marginals $= (1 + \epsilon/2)z_e$ or the ellipsoid method (see [2] section 8) where the basic idea is to find a near optimal solution of the dual of the convex program.

3 Inapproximability Bounds of ATSP

4 Special Cases of ATSP

4.1 ATSP on Graphs with Bounded Genus

4.2 Polyloglogn Integrality Gap of ATSP

4.3 Local-Connectivity ATSP

5 Concluding Remarks

References

- [1] ANARI, N., AND GHARAN, S. O. Effective-resistance-reducing flows, spectrally thin trees, and asymmetric TSP. *In: Proceedings of FOCS* (2015).
- [2] ASADPOUR, A., GOEMANS, M. X., MADRY, A., GHARAN, S. O., AND SABERI, A. An $\mathcal{O}(\log n / \log \log n)$ -approximation algorithm for the asymmetric traveling salesman problem. *In: Proceedings of SODA* (2010), pages 379–389.
- [3] CHRISTOFIDES, N. Worst-case analysis of a new heuristic for the traveling salesman problem. *Technical Report, DTIC Document* (1976).
- [4] EDMONDS, J. Matroids and the greedy algorithm. *Mathematical Programming* (1971), pages 127–136.
- [5] FRIEZE, A. M., GALBIATI, G., AND MAFFIOLI, F. On the worse-case performance of some algorithms for the asymmetric traveling salesman problem. *Networks* (1982), 12:23–39.
- [6] GHARAN, S. O., AND SABERI, A. The asymmetric traveling salesman problem on graphs with bounded genus. *In: Proceedings of SODA* (2011), pages 967–975.
- [7] GHARAN, S. O., SABERI, A., AND SINGH, M. A randomized rounding approach to the traveling salesman problem. *In: Proceedings of FOCS* (2011).
- [8] GOEMANS, M. X. Minimum bounded degree spanning trees. *In FOCS* (2006), pages 273–282.
- [9] HELD, M., AND KARP, R. The traveling salesman problem and minimum spanning trees. *Operations Research* (1970), 18:1138–1162.
- [10] KARPINSKI, M., LAMPIS, M., AND SCHMIED, R. New approximability bounds for TSP. *J. Comput. Syst. Sci.* (2015), 81(8):1665–1667.

- [11] PAPADIMITRIOU, C. H., AND YANNAKAKIS, M. The traveling salesman problem with distances one and two. *Math. Oper. Res.* (1993), 18(1):1–11.
- [12] SCHRIJVER, A. Combinatorial optimization. *vol. 2 of Algorithms and Combinatorics* (2003).
- [13] SVENSSON, O. Approximating ATSP by relaxing connectivity. *In: Proceedings of FOCS* (2015).
- [14] SVENSSON, O., TARNAWSKI, J., AND VÉGH, L. A. Constant factor approximation for ATSP with two edge weights. *In: Proceedings of IPCO* (2016), pages 226–237.